

Student name: Gordon Tse

Student ID:s223532749

# SIT225: Data Capture Technologies

## Activity 7.1: Data analysis and interpretation

Data analysis is a broad term that covers a wide range of techniques that enable you to reveal any insights and relationships that may exist within raw data. As you might expect, Python lends itself readily to data analysis. Once Python has analyzed your data, you can then use your findings to make good business decisions, improve procedures, and even make informed predictions based on what you've discovered.

You have done data wrangling using Python Pandas module already in activity 5.2. In this activity, you will learn Data science statistics and linear regression models.

### Hardware Required

No hardware is required.

### Software Required

Python 3



Python packages including Pandas, Numpy, Scikit-learn, seaborn, plotly

### Steps:

| Step | Action                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
|------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1    | A Jupyter Notebook is provided for Data Science exploration here ( <a href="https://github.com/deakin-deep-dreamer/sit225/tree/main/week_7">https://github.com/deakin-deep-dreamer/sit225/tree/main/week_7</a> ). You will need to fill in your student ID and name and run all the cells to observe the output. Convert the Notebook into PDF and merge with this activity sheet which needs to be combined with this week's task for OnTrack submission. |

|   |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
|---|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|   | <p><b>Question:</b> There are sections in the Notebook. After running the cells and observing the outputs, provide your reflection in brief on the topic items for each section of the Notebook.</p> <p><b>Answer:</b></p> <p><b>1. Descriptive Statistics</b> – this section I learn about different math statistics like sum, avg, count and percentile.</p> <p><b>1.1 Percentile</b> – it let me find the what is the number at different %.</p> <p><b>1.2 Standard Deviation</b> – it showing me the core topic of Statistics, this show how spread are the data. <b>Coefficient of variation</b> show how large is the sdt is.</p> <p><b>1.3 Variance</b> - multiply the standard deviation by itself = Variance. It also a way to show how spread are the data.</p> <p><b>1.4 Correlation</b> – to show the relationship between two variables. It show if x change, what will happened to y.<br/> 1 = perfect linear relationship: x increase y increase.<br/> 0 = no linear relationship: x increase y no change.<br/> -1 = perfect negative linear relationship: x increase y decrease.</p> <p><b>1.5 Correlation Matrix</b> – a array to show correlation.</p> <p><b>1.6 Correlation Does not imply Causality</b> – what I understand is when correlation cannot be apply to the variables. Like they have no relationship But getting Correlation = 1 or -1.</p> <p><b>1.7 Linear Regression</b> – it help to draw a straight line to show the relationship data. It can use for predict.</p> |
| 2 | <p><b>Question:</b> In the 1.1 Percentile subsection of <b>Descriptive statistics</b> section in the Notebook, you have calculated 10%, 25%, 50% and 75% percentiles for <i>Max_Pulse</i>. Compare these percentiles with <i>Average_Pulse</i> percentiles for any trend, if exists.</p> <p><b>Answer:</b> now I get the percentile for max_pulse, I found that overall in max_pulse is more then average_pulse which is normal because it showing the max number. also both when percentile increase the number also increase .</p>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |

|   |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
|---|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|   | <pre> avg_pulse = full_health_data["Average_Pulse"] print("percentile_10", np.percentile(avg_pulse, 10) ) print("percentile_25", np.percentile(avg_pulse, 25) ) print("percentile_50", np.percentile(avg_pulse, 50) ) print("percentile_75", np.percentile(avg_pulse, 75) )  ✓ 0.0s  percentile_10 92.2 percentile_25 100.0 percentile_50 105.0 percentile_75 111.0  max_pulse = full_health_data["Max_Pulse"] print("percentile_10", np.percentile(max_pulse, 10) ) print("percentile_25", np.percentile(max_pulse, 25) ) print("percentile_50", np.percentile(max_pulse, 50) ) print("percentile_75", np.percentile(max_pulse, 75) )  ✓ 0.0s  percentile_10 120.0 percentile_25 124.0 percentile_50 131.0 percentile_75 141.0 </pre>                                                              |
| 3 | <p><b>Question:</b> In the “Correlation Does not imply Causality” section answer the question regarding the increase of ice cream sale in your own understanding.</p> <p><b>Answer:</b> Correlation is the measures the relation between two variables, when we talk about during summer both sale of ice cream and drowning accidents increase, I will say it kind of having Correlation because of warm weather. The weather is the cause. But since when weather getting hot, people will buy ice cream and go to swimming, so I will they have the correlation. <b>In short the direct cause of both event is warm weather.</b></p>                                                                                                                                                             |
| 4 | <p><b>Question:</b> In the <b>1.7 Linear Regression</b> section in the Notebook, a linear regression model was used to predict Calorie_Burnage from attributes such as Average_Pulse. The Duration value was predicted from the model for all the value range of Average_Pulse and a regression line was drawn. You will need to answer the follow up question next to 1.7 section where it is required to generate a linear regression model for Duration instead of Average_Pulse to predict the Calorie_Burnage. Take a screenshot of the regression line and paste it here. Also, comment on both the regression lines.</p> <p><b>Answer:</b></p> <p>Looking at the Calorie_Burnage – Duration graph, is showing a upward trend and Calorie_Burnage - Average_Pulse is showing a flat line.</p> |





# SIT225: Data Analysis & interpretation

Run each cell to generate output and finally convert this notebook to PDF.

```
In [1]: # Fill in student ID and name
#
student_id = "s223532749"
student_first_last_name = "Gordon Tse"
print(student_id, student_first_last_name)
```

s223532749 Gordon Tse

## 1. Descriptive Statistics

Descriptive statistics summarizes important features of a data set such as:

- Count
- Sum
- Standard deviation
- Percentile
- Average

```
In [2]: # Make sure necessary packages are already installed.
# !pip install pandas numpy seaborn

import pandas as pd
import numpy as np
import seaborn as sns

full_health_data = pd.read_csv("full_health_data.csv", header=0, sep=",")
print (full_health_data.describe())
```

|       | Duration   | Average_Pulse | Max_Pulse  | Calorie_Burnage | Hours_Work \ |
|-------|------------|---------------|------------|-----------------|--------------|
| count | 163.000000 | 163.000000    | 163.000000 | 163.000000      | 163.000000   |
| mean  | 64.263804  | 107.723926    | 134.226994 | 382.368098      | 4.386503     |
| std   | 42.994520  | 14.625062     | 16.403967  | 274.227106      | 3.923772     |
| min   | 15.000000  | 80.000000     | 100.000000 | 50.000000       | 0.000000     |
| 25%   | 45.000000  | 100.000000    | 124.000000 | 256.500000      | 0.000000     |
| 50%   | 60.000000  | 105.000000    | 131.000000 | 320.000000      | 5.000000     |
| 75%   | 60.000000  | 111.000000    | 141.000000 | 388.500000      | 8.000000     |
| max   | 300.000000 | 159.000000    | 184.000000 | 1860.000000     | 11.000000    |

|       | Hours_Sleep |
|-------|-------------|
| count | 163.000000  |
| mean  | 7.680982    |
| std   | 0.663934    |
| min   | 5.000000    |
| 25%   | 7.500000    |
| 50%   | 8.000000    |
| 75%   | 8.000000    |
| max   | 12.000000   |

## 1.1 Percentile

### 25%, 50% and 75% - Percentiles

Observe the output of the above cell for 25%, 50% and 75% of all the columns.

Let's explain for Average\_Pulse:

- 25% of all of the training sessions have an average pulse of 100 beats per minute or lower. If we flip the statement, it means that 75% of all of the training sessions have an average pulse of 100 beats per minute or higher.
- 75% of all the training session have an average pulse of 111 or lower. If we flip the statement, it means that 25% of all of the training sessions have an average pulse of 111 beats per minute or higher.

```
In [3]: avg_pulse = full_health_data["Average_Pulse"]
print("percentile_10", np.percentile(avg_pulse, 10) )
print("percentile_25", np.percentile(avg_pulse, 25) )
print("percentile_50", np.percentile(avg_pulse, 50) )
print("percentile_75", np.percentile(avg_pulse, 75) )
```

```
percentile_10 92.2
percentile_25 100.0
percentile_50 105.0
percentile_75 111.0
```

**Question:** Calculate percentiles for Max\_Pulse.

You should answer a follow up question in the activity sheet.

```
In [4]: max_pulse = full_health_data["Max_Pulse"]
print("percentile_10", np.percentile(max_pulse, 10) )
print("percentile_25", np.percentile(max_pulse, 25) )
print("percentile_50", np.percentile(max_pulse, 50) )
print("percentile_75", np.percentile(max_pulse, 75) )
```

```
percentile_10 120.0
percentile_25 124.0
percentile_50 131.0
percentile_75 141.0
```

## 1.2 Standard Deviation

Standard deviation is a number that describes how spread out the observations are.

A mathematical function will have difficulties in predicting precise values, if the observations are "spread". Standard deviation is a measure of uncertainty.

A low standard deviation means that most of the numbers are close to the mean (average) value.

A high standard deviation means that the values are spread out over a wider range.

```
In [5]: import numpy as np

# We can use the std() function from Numpy to find the standard deviation of a

std = np.std(full_health_data)
print(std)
```

```
Duration          42.862432
Average_Pulse     14.580131
Max_Pulse         16.353571
Calorie_Burnage   273.384624
Hours_Work        3.911718
Hours_Sleep       0.661895
dtype: float64
```

```
c:\Users\gordo\anaconda3\Lib\site-packages\numpy\core\fromnumeric.py:3643: FutureWarning: The behavior of DataFrame.std with axis=None is deprecated, in a future version this will reduce over both axes and return a scalar. To retain the old behavior, pass axis=0 (or do not pass axis)
    return std(axis=axis, dtype=dtype, out=out, ddof=ddof, **kwargs)
```

### 1.2.1 Coefficient of variation

In the above cell, what does standard deviation numbers mean?

The coefficient of variation is used to get an idea of how large the standard deviation is.

Mathematically, the coefficient of variation is defined as:  $\text{Coefficient of Variation} = \text{Standard Deviation} / \text{Mean}$

```
In [6]: cv = np.std(full_health_data) / np.mean(full_health_data)
print(cv)

# We see that the variables Duration and Calorie_Burnage has
# a high Standard Deviation compared to Max_Pulse, Average_Pulse and Hours_Sleep
#
```

```
Duration          0.367051
Average_Pulse     0.124857
Max_Pulse         0.140043
Calorie_Burnage   2.341122
Hours_Work        0.033498
Hours_Sleep       0.005668
dtype: float64
```

## 1.3 Variance

Variance is another number that indicates how spread out the values are.

In fact, if you take the square root of the variance, you get the standard deviation. Or the other way around, if you multiply the standard deviation by itself, you get the variance!

```
In [7]: var = np.var(full_health_data)
print(var)
```

```
Duration          1837.188076
Average_Pulse     212.580225
Max_Pulse         267.439271
Calorie_Burnage   74739.152847
Hours_Work        15.301536
Hours_Sleep       0.438105
dtype: float64
```

```
c:\Users\gordo\anaconda3\Lib\site-packages\numpy\core\fromnumeric.py:3785: FutureWarning: The behavior of DataFrame.var with axis=None is deprecated, in a future version this will reduce over both axes and return a scalar. To retain the old behavior, pass axis=0 (or do not pass axis)
  return var(axis=axis, dtype=dtype, out=out, ddof=ddof, **kwargs)
```

## 1.4 Correlation

Correlation measures the relationship between two variables.

A function has a purpose to predict a value, by converting input (x) to output (f(x)). We can say also that a function uses the relationship between two variables for prediction.

## Correlation Coefficient

The correlation coefficient measures the relationship between two variables.

The correlation coefficient can never be less than -1 or higher than 1.

- 1 = there is a perfect linear relationship between the variables
- 0 = there is no linear relationship between the variables
- -1 = there is a perfect negative linear relationship between the variables

### Perfect Linear Relationship (Correlation Coefficient = 1)


it exists a perfect linear relationship between Average\_Pulse and Calorie\_Burnage.

```
In [8]: # Positive correlation
#

import matplotlib.pyplot as plt

def create_linear_health_data():
    data = [
        {'Duration': 30, 'Average_Pulse': 80, 'Max_Pulse': 120, 'Calorie_Burnage': 100},
        {'Duration': 45, 'Average_Pulse': 85, 'Max_Pulse': 120, 'Calorie_Burnage': 150},
        {'Duration': 45, 'Average_Pulse': 90, 'Max_Pulse': 130, 'Calorie_Burnage': 200},
        {'Duration': 60, 'Average_Pulse': 95, 'Max_Pulse': 130, 'Calorie_Burnage': 250},
        {'Duration': 60, 'Average_Pulse': 100, 'Max_Pulse': 140, 'Calorie_Burnage': 300},
        {'Duration': 60, 'Average_Pulse': 105, 'Max_Pulse': 140, 'Calorie_Burnage': 350},
        {'Duration': 60, 'Average_Pulse': 110, 'Max_Pulse': 145, 'Calorie_Burnage': 400},
        {'Duration': 45, 'Average_Pulse': 115, 'Max_Pulse': 145, 'Calorie_Burnage': 450},
        {'Duration': 60, 'Average_Pulse': 120, 'Max_Pulse': 150, 'Calorie_Burnage': 500},
        {'Duration': 45, 'Average_Pulse': 125, 'Max_Pulse': 150, 'Calorie_Burnage': 550}
    ]
    return data

health_data = pd.DataFrame.from_dict(create_linear_health_data())
health_data.plot(x='Average_Pulse', y='Calorie_Burnage', kind='scatter')
plt.show()
```



### Perfect Negative Linear Relationship (Correlation Coefficient = -1)


We have plotted fictional data here. The x-axis represents the amount of hours worked at our job before a training session. The y-axis is Calorie\_Burnage.

If we work longer hours, we tend to have lower calorie burnage because we are exhausted before the training session.

The correlation coefficient here is -1.

```
In [9]: # Negative correlation
#
negative_corr = {'Hours_Work_Before_Training': [10,9,8,7,6,5,4,3,2,1],
                 'Calorie_Burnage': [220,240,260,280,300,320,340,360,380,400]}
negative_corr = pd.DataFrame(data=negative_corr)

negative_corr.plot(x='Hours_Work_Before_Training', y='Calorie_Burnage', kind=
plt.show())
```




No Linear Relationship (Correlation coefficient = 0)

As you can see, there is no linear relationship between the two variables. It means that longer training session does not lead to higher Max\_Pulse.

The correlation coefficient here is 0.

```
In [10]: full_health_data.plot(x='Duration', y='Max_Pulse', kind='scatter')
plt.show()
```



## 1.5 Correlation Matrix

A matrix is an array of numbers arranged in rows and columns.

A correlation matrix is simply a table showing the correlation coefficients between variables.

We can use the `corr()` function in Python to create a correlation matrix. We also use the `round()` function to round the output to two decimals:

```
In [11]: Corr_Matrix = round(full_health_data.corr(),2)
print(Corr_Matrix)

# Drop 2 columns - Hours_Work and Hours_Sleep to view the matrix nice.
#
health_part = full_health_data.drop(columns=['Hours_Work', 'Hours_Sleep'])
Corr_Matrix = round(health_part.corr(),2)
print(Corr_Matrix)
```

|                 | Duration | Average_Pulse | Max_Pulse | Calorie_Burnage | \ |
|-----------------|----------|---------------|-----------|-----------------|---|
| Duration        | 1.00     | -0.17         | 0.00      | 0.89            |   |
| Average_Pulse   | -0.17    | 1.00          | 0.79      | 0.02            |   |
| Max_Pulse       | 0.00     | 0.79          | 1.00      | 0.20            |   |
| Calorie_Burnage | 0.89     | 0.02          | 0.20      | 1.00            |   |
| Hours_Work      | -0.12    | -0.28         | -0.27     | -0.14           |   |
| Hours_Sleep     | 0.07     | 0.03          | 0.09      | 0.08            |   |

|                 | Hours_Work | Hours_Sleep |
|-----------------|------------|-------------|
| Duration        | -0.12      | 0.07        |
| Average_Pulse   | -0.28      | 0.03        |
| Max_Pulse       | -0.27      | 0.09        |
| Calorie_Burnage | -0.14      | 0.08        |
| Hours_Work      | 1.00       | -0.14       |
| Hours_Sleep     | -0.14      | 1.00        |

|                 | Duration | Average_Pulse | Max_Pulse | Calorie_Burnage |
|-----------------|----------|---------------|-----------|-----------------|
| Duration        | 1.00     | -0.17         | 0.00      | 0.89            |
| Average_Pulse   | -0.17    | 1.00          | 0.79      | 0.02            |
| Max_Pulse       | 0.00     | 0.79          | 1.00      | 0.20            |
| Calorie_Burnage | 0.89     | 0.02          | 0.20      | 1.00            |

Using a Heatmap


We can use a Heatmap to Visualize the Correlation Between Variables:

```
In [12]: import matplotlib.pyplot as plt
import seaborn as sns

correlation_full_health = full_health_data.corr()

axis_corr = sns.heatmap(
    correlation_full_health,
    vmin=-1, vmax=1, center=0,
    cmap=sns.diverging_palette(50, 500, n=500),
    square=True
)

plt.show()
```



## 1.6 Correlation Does not imply Causality

Correlation measures the numerical relationship between two variables.

A high correlation coefficient (close to 1), does not mean that we can for sure conclude an actual relationship between two variables.

A classic example:

- During the summer, the sale of ice cream at a beach increases
- Simultaneously, drowning accidents also increase as well

**Question:** Does this mean that increase of ice cream sale is a direct cause of increased drowning accidents?

## 1.7 Linear Regression

The term regression is used when you try to find the relationship between variables.

In Machine Learning and in statistical modeling, that relationship is used to predict the outcome of events.

We will use Scikit-learn to train various regression models. Scikit-learn is a popular Machine Learning (ML) library that offers various tools for creating and training ML algorithms, feature engineering, data cleaning, and evaluating and testing models. It was designed to be accessible, and to work seamlessly with popular libraries like NumPy and Pandas.

We see how to apply a simple regression model for predicting Calorie\_Burnage on various factors such as Average\_Pulse or Duration.

```
In [13]: # !pip install seaborn plotly

import numpy as np
import plotly.express as px
import plotly.graph_objects as go
from sklearn.linear_model import LinearRegression

df = full_health_data
X = df.Average_Pulse.values.reshape(-1, 1)

model = LinearRegression()
model.fit(X, df.Calorie_Burnage)

x_range = np.linspace(X.min(), X.max(), 100)
y_range = model.predict(x_range.reshape(-1, 1))

fig = px.scatter(df, x='Average_Pulse', y='Calorie_Burnage', opacity=0.65)
fig.add_traces(go.Scatter(x=x_range, y=y_range, name='Regression Fit'))
fig.show()
```

### Question:

We have seen earlier how to apply a simple regression model for predicting Calorie\_Burnage from Average\_Pulse. There might be another candidate Duration in addition to Average\_Pulse. You will need to repeat the above linear regression process to find relationship between Calorie\_Burnage and Duration.

Comment on the both regression lines: Calorie\_Burnage - Average\_Pulse and Calorie\_Burnage - Duration.

```
In [15]: import numpy as np
import plotly.express as px
import plotly.graph_objects as go
from sklearn.linear_model import LinearRegression

df = full_health_data
X = df.Duration.values.reshape(-1, 1)

model = LinearRegression()
model.fit(X, df.Calorie_Burnage)

x_range = np.linspace(X.min(), X.max(), 100)
y_range = model.predict(x_range.reshape(-1, 1))

fig = px.scatter(df, x='Duration', y='Calorie_Burnage', opacity=0.65)
fig.add_traces(go.Scatter(x=x_range, y=y_range, name='Regression Fit'))
fig.show()
```